

```
In [4]: # Step 1: Handling Missing Values
imputer = SimpleImputer(strategy='mean')
X_imputed = imputer.fit_transform(X)
```

```
In [5]: X_imputed

Out[5]: array([[1.799e+01, 1.038e+01, 1.228e+02, ..., 2.654e-01, 4.601e-01,
1.189e-01],
[2.057e+01, 1.777e+01, 1.329e+02, ..., 1.860e-01, 2.750e-01,
8.902e-02],
[1.969e+01, 2.125e+01, 1.300e+02, ..., 2.430e-01, 3.613e-01,
8.758e-02],
...,
[1.660e+01, 2.808e+01, 1.083e+02, ..., 1.418e-01, 2.218e-01,
7.820e-02],
[2.060e+01, 2.933e+01, 1.401e+02, ..., 2.650e-01, 4.087e-01,
1.240e-01],
[7.760e+00, 2.454e+01, 4.792e+01, ..., 0.000e+00, 2.871e-01,
7.039e-02]])

In [6]: # Step 3: Normalization or Scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_imputed)

In [7]: # Step 4: Feature Engineering
# Introduce an artificial time column
time_column = np.arange(X_scaled.shape[0]).reshape(-1, 1)
X_with_time = np.hstack((X_scaled, time_column))

In [8]: # Create time-based features
X_with_time = pd.DataFrame(X_with_time, columns=[f'feature_{i}' for i in range(X_scaled.shape[1])] + ['time'])
X_with_time['month'] = X_with_time['time'] % 12 + 1
X_with_time['day_of_week'] = X_with_time['time'] % 7 + 1

In [9]: # Rolling statistics
rolling_window = 5
for i in range(X_scaled.shape[1]):
    X_with_time[f'rolling_mean_{i}'] = X_with_time[f'feature_{i}'].rolling(window=rolling_window).mean()
    X_with_time[f'rolling_std_{i}'] = X_with_time[f'feature_{i}'].rolling(window=rolling_window).std()

In [10]: X_with_time

Out[10]:
   feature_0 feature_1 feature_2 feature_3 feature_4 feature_5 feature_6 feature_7 feature_8 feature_9 ... rolling_mean_25 rolling_std_25 rolling_mean_26 rolling_std_26 rolling_mean_27 rolling_std_27
0    1.097064 -2.074303  1.270930  0.984375  1.568466  3.283515  2.652874  2.532475  2.217515  2.255747 ...              NaN              NaN              NaN              NaN              NaN              NaN
1    1.829821 -0.351816  1.687023  1.908708 -0.826962 -0.487072 -0.023846  0.548144  0.001392 -0.868652 ...              NaN              NaN              NaN              NaN              NaN              NaN
2    1.579888  0.459314  1.567551  1.558884  0.942210  1.052926  1.363478  2.037231  0.939685 -0.398008 ...              NaN              NaN              NaN              NaN              NaN              NaN
3   -0.768909  0.256531 -0.592011 -0.764464  3.283553  3.402909  1.915897  1.451707  2.867383  4.910919 ...              NaN              NaN              NaN              NaN              NaN              NaN
4    1.750297 -1.151293  1.777657  1.826229  0.280372  0.539340  1.371011  1.428493 -0.009560 -0.562450 ...    1.369831    1.876135    1.084104    0.956643    1.648641    0.698448
...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...
564  2.110995  0.725029  2.061918  2.343856  1.041842  0.219060  1.947285  2.320965 -0.312589 -0.931027 ...    0.549135    1.801907    0.971266    2.229958    0.740527    1.650113
565  1.704854  2.090898  1.616988  1.723842  0.102458 -0.017833  0.693043  1.263669 -0.217664 -1.058611 ...    0.505623    1.826295    1.152516    2.096048    0.917155    1.576711
566  0.702284  2.051274  0.673570  0.577953 -0.840484 -0.038680  0.046588  0.105777 -0.809117 -0.895587 ...    0.829367    1.556893    1.479036    1.708726    1.348982    0.737801
567  1.838341  2.342628  1.983644  1.735218  1.525767  3.272144  3.296944  2.658866  2.137194  1.043695 ...    0.926569    1.760799    1.257102    1.263652    1.438514    0.833558
568 -1.808401  1.226158 -1.813922 -1.347789 -3.112085 -1.150752 -1.114873 -1.261820 -0.820070 -0.561032 ...    0.475979    1.995422    0.623925    1.626943    0.664394    1.536731

569 rows x 93 columns
```

```
In [11]: # Additional transformations
X_with_time['log_feature_0'] = np.log1p(X_with_time['feature_0'])
X_with_time['sqrt_feature_1'] = np.sqrt(X_with_time['feature_1'])

C:\Users\Chiranjeevi\anaconda3\lib\site-packages\pandas\core\arraylike.py:358: RuntimeWarning: invalid value encountered in log1p
    result = getattr(ufunc, method)(*inputs, **kwargs)
C:\Users\Chiranjeevi\anaconda3\lib\site-packages\pandas\core\arraylike.py:358: RuntimeWarning: invalid value encountered in sqrt
    result = getattr(ufunc, method)(*inputs, **kwargs)

In [12]: # Drop rows with missing values introduced during rolling statistics computation
X_with_time = X_with_time.dropna()

In [13]: X_with_time

Out[13]:
   feature_0 feature_1 feature_2 feature_3 feature_4 feature_5 feature_6 feature_7 feature_8 feature_9 ... rolling_mean_26 rolling_std_26 rolling_mean_27 rolling_std_27 rolling_mean_28 rolling_s
6    1.170908  0.163298  1.139099e+00  1.095295 -0.123136  0.088295  0.300072  0.646935 -0.064325 -0.762332 ...    1.046107    0.601904    1.392530    0.641279    1.669298    2.61
7   -0.118517  0.361419  -7.210105e-02 -0.218965  1.604049  1.140102  0.061026  0.281950  1.403355  1.660353 ...    0.870902    0.774269    1.126369    0.625442    1.534375    2.61
8   -0.320167  0.592171  -1.833337e-01 -0.384207  2.201839  1.684010  1.219096  1.150692  1.965600  1.572462 ...    0.728991    0.550728    0.969535    0.320471    0.803138    1.21
9   -0.473535  1.109617  -3.287601e-01 -0.509063  1.582699  2.563358  1.738872  0.941760  0.797298  2.783096 ...    1.405442    1.547701    1.147686    0.392909    1.450897    1.01
10  0.537556  0.923150  5.854486e-16  0.406453 -1.017686 -0.713542 -0.700684 -0.404686 -1.035476 -0.826124 ...    1.031619    1.796382    0.921267    0.739984    1.115370    1.11
...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...
563  1.929226  1.354354  2.103116e+00  1.968434  0.963560  2.260135  2.870075  2.540213  1.231760  0.849484 ...    0.925498    2.240153    0.359649    1.613120    -0.504029    1.51
564  2.110995  0.725029  2.061918e+00  2.343856  1.041842  0.219060  1.947285  2.320965 -0.312589 -0.931027 ...    0.971266    2.229958    0.740527    1.650139    -0.520854    1.51
565  1.704854  2.090898  1.616988e+00  1.723842  0.102458 -0.017833  0.693043  1.263669 -0.217664 -1.058611 ...    1.152516    2.096048    0.917155    1.576712    -0.416669    1.51
566  0.702284  2.051274  6.735698e-01  0.577953 -0.840484 -0.038680  0.046588  0.105777 -0.809117 -0.895587 ...    1.479036    1.708726    1.348982    0.737801    -0.205710    1.31
567  1.838341  2.342628  1.983644e+00  1.735218  1.525767  3.272144  3.296944  2.658866  2.137194  1.043695 ...    1.257102    1.263652    1.438514    0.833580    -0.206357    1.31

239 rows x 95 columns
```

```
In [14]: # Drop rows with missing values introduced during rolling statistics computation
X_with_time = X_with_time.dropna()

# Update the target variable 'y' accordingly
y = y[:X_with_time.shape[0]]

# Split the Data
X_train, X_test, y_train, y_test = train_test_split(X_with_time.drop(['time'], axis=1), y, test_size=0.2, random_state=42)
```

In [15]:

X_train

Out[15]:

	feature_0	feature_1	feature_2	feature_3	feature_4	feature_5	feature_6	feature_7	feature_8	feature_9	...	rolling_mean_26	rolling_std_26	rolling_mean_27	rolling_std_27	rolling_mean_28	rolling_std_2
302	1.693494	1.067662	1.761178	1.684024	0.828346	1.505866	1.751427	2.039810	1.596855	1.685870	...	0.127458	1.160639	0.000210	1.153491	-0.224476	0.64381
119	1.085703	0.170290	0.916634	0.930337	-0.878202	-0.703498	-0.199239	0.181612	1.158741	-1.778754	...	0.359486	1.215894	0.190329	1.140789	0.429431	1.97376
506	-0.541698	0.177283	-0.514148	-0.573909	0.942210	0.205794	-0.088504	-0.703122	1.140487	0.870748	...	0.059545	0.405083	0.086362	0.919967	0.049899	0.27824
413	0.245021	0.659766	0.229875	0.110382	-0.797785	-0.034889	-0.253727	-0.262045	0.483317	-0.519922	...	-0.419536	0.374416	-0.487194	0.580680	0.149231	0.63869
503	2.545536	0.128335	2.478011	2.921209	-0.209246	0.438898	0.989340	1.325317	-1.119447	-1.128074	...	0.238876	0.671016	0.617222	0.966967	-0.042961	0.93106
...
232	-0.825712	3.386842	-0.871740	-0.762473	-1.320851	-1.300090	-1.052512	-1.095861	0.121873	-0.640418	...	0.280403	1.235289	0.145624	1.558811	0.193558	0.39333
26	0.128576	0.524577	0.224931	-0.028694	0.643316	1.562720	0.674211	1.003666	1.607807	0.913276	...	0.953515	0.606083	1.832579	0.448057	1.594556	1.14161
201	0.969258	0.009463	0.953711	0.843876	-0.475405	0.292971	0.185822	0.669634	-1.115796	-1.118151	...	0.253556	0.653073	0.533780	0.708151	0.455315	1.50415
446	1.028901	2.039619	1.044345	0.929199	0.256887	0.512808	1.016961	0.877275	-0.360051	-0.515669	...	-0.053840	1.273178	-0.143103	1.156278	-0.495940	0.70315
223	0.460872	0.226230	0.438745	0.302644	0.436936	0.304342	0.325182	0.404988	0.450458	0.032944	...	-0.061910	0.691968	-0.107869	0.695130	0.227208	0.97456

191 rows × 94 columns

In [16]:

```
# Step 6: Pipeline with PCA and RandomForestClassifier
pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy='mean')),
    ('scaler', StandardScaler()),
    ('pca', PCA(n_components=10)),
    ('classifier', RandomForestClassifier(random_state=42))
])
```

In [17]:

```
# Fit the model
pipeline.fit(X_train, y_train)
```

Out[17]:

```
graph TD
    subgraph Pipeline
        direction TB
        A[SimpleImputer] --> B[StandardScaler]
        B --> C[PCA]
        C --> D[RandomForestClassifier]
    end
```

In [18]:

```
# Step 8: Monitor Performance Metrics
y_pred = pipeline.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
```

Accuracy: 0.42

In []: